

Guide to the MLR Code

M Toussaint

November 30, 2016

Contents

1	Graph	1
2	Logic	3
2.1	Methods	4
3	Optim	4
3.1	Representing Optimization Problems	5
4	KOMO	6
4.1	Formal problem representation	7
4.2	User Interfaces	8
4.2.1	Easy	8
4.2.2	Using a specs file	9
4.2.3	Expert using the included kinematics engine	9
4.2.4	Expert with own kinematics engine	9
4.2.5	Optimizers	10
4.2.6	Parameters & Reporting	10
4.3	Potential Improvements	10
4.4	Disclaimer	10
5	Control	10
5.1	Slow and fast control loop	10
5.2	Operational Space Control: Computing gains by projecting operational space gains	11
5.3	Controlling the F/T signal—the <i>sensed</i> force	12
5.3.1	Preliminaries: Understanding force transmission	12
5.3.2	Controlling the direct F/T signal with a fixated endeff	12
5.3.3	Control the indirectly sensed contact force of endeff	13
5.3.4	q -control under force constraints	14
5.3.5	I-gains on position?	14
5.4	Technical Details and Issues	14
5.4.1	Ctrl-Message documentation	14
5.4.2	Filtering of the differentiation of q	14
5.5	Enforcing control, velocity, joint and force limits	15
.1	Reference: General OPEN-LOOP ideal contact force controller	15
.1.1	General case	15
.2	Reference: Pullback of operational space linear controllers	16
.3	How to make this FEEDBACK?	16
.4	What do we want?	16

1 Graph

Our graph syntax is a bit different to standard conventions. Actually, our graph could be called a *key-value hierarchical hyper graph*: nodes can play the role of normal nodes, or hypernodes (=edges or factors/cliques) that connect other nodes. Every node also has a set

of keys (or tags, to retrieve the node by name) and a typed value (every node can be of a different type). This value can also be a graph, allowing to represent hierarchies of graphs and subgraphs.

- A graph is a set of nodes
- Every node has three properties:
 - A tuple of **keys** (=strings)
 - A tuple of **parents** (=references to other nodes)
 - A typed **value** (the type may differ for every node)

Therefore, depending on the use case, such a graph could represent just a key-value list, an 'any-type' container (container of things of varying types), a normal graph, a hierarchical graph, or an xml data structure.

We use the graph in particular also to define a generic file format, which we use for configuration (parameter) files, files that define robot kinematic and geometry, or any other structured data. This ascii file format of a graph helps to also understand the data structure. Here is the `example.g` from `test/Core/graph`:

```
## a trivial graph
x      # a 'vertex': key=x, value=true, parents=None
y      # another 'vertex': key=y, value=true, parents=None
(x y)  # an 'edge': key=None, value=true, parents=x y
x      # key=x, value=true, parents=None
y      # key=y, value=true, parents=None
(x y)  # key=None, value=true, parents=x y
(-1 -2) # key=None, value=true, parents=the previous and the y-node

## a bit more verbose graph
node A{ color=blue } # keys=node A, value=<Graph>, parents=None
node B{ color=red, value=5 } # keys=node B, value=<Graph>, parents=None
edge C(A,B){ width=2 } # keys=edge C, value=<Graph>, parents=A B
hyperedge(A B C) = 5 # keys=hyperedge, value=5, parents=A B C

## standard value types
a=string # MT::String (except for keywords 'true' and 'false' and 'Mod' and 'Include')
b="STRING" # MT::String (does not require a '=')
c='file.txt' # MT::FileToken (does not require a '=')
d=-0.1234 # double
e=[1 2 3 0.5] # MT::arr (does not require a '=')
#f=(c d e) # DEPRECATED!! MT::Array<Node> (list of other nodes in the Graph)
g! # bool (default: true, !means false)
h=true # bool
i=false # bool
j={ a=0 } # sub-Graph (special: does not require a '=')

## parsing: = {...} (...), and \n are separators for parsing key-value-pairs
b0=false b1 b2, b3() b4 # 4 boolean nodes with keys 'b0', 'b1 b2', 'b3', 'b4'
k=( a, b=0.2 x="hallo" # sub-Graph with 6 nodes
  y
  z()=filename.org x )

## special Node Keys

## editing: after reading all nodes, the Graph takes all Edit nodes, deletes the Edit tag, and calls a edit()
## this example will modify/append the respective attributes of k
Edit k { y=false, z=otherString, b=7, c=newAttribute }

## including
Include = 'example_include.g' # first creates a normal FileToken node then opens and includes the file directly

## any types
#trans=<T t(10 0 0)> # 'T' is the tag for an arbitrary type (here an ors::Transformation)
# which was registered somewhere in the code using the registry()
# (does not require a '=')

## strange notations
a() # key=a, value=true, parents=None
() # key=None, value=true, parents=None
[1 2 3 4] # key=None, value=MT::arr, parents=None
```

A special case is when a node has a Graph-typed value. This is considered a **subgraph**. Subgraphs are sometimes handled special: their nodes can have parents from the containing graph, or from other subgraphs of the containing graph. Some methods of the `Graph`

class (to find nodes by key or type) allow to specify whether also nodes in subgraphs or parentgraphs are to be searched.

2 Logic

We represent everything, a full knowledge base (KB), as a graph:

- Symbols (both, constants and predicate symbols) are nil-valued nodes. We assume that they are declared in the root scope of the graph
- A grounded literal is a tuple of symbols, for instance `(on box1 box2)`. Note that we can equally write this as `(box1 on box2)`. There is no need to have the ‘predicate’ first. In fact, the basic methods do not distinguish between predicate and contant symbols.
- A universal quantification $\forall X$ is represented as a scope (=subgraph) which first declares the logic variables as nil-valued nodes as the subgraph, then the rest. The rest is typically an implication, i.e., a rule. For instance

$$\forall XY \ p(X, Y)q(Y) \Rightarrow q(X)$$

is represented as `{X, Y, { (p X Y) (q Y) } { (q X) } }` where the precondition and postconditions are subgraphs of the rule-subgraph.

Here is how the standard FOL example from Stuart Russell’s lecture is represented:

```

Constant M1
Constant M2
Constant Nono
Constant America
Constant West

American
Weapon
Sells
Hostile
Criminal
Missile
Owns
Enemy

STATE {
  (Owns Nono M1),
  (Missile M1),
  (Missile M2),
  (American West),
  (Enemy Nono America)
}

Query { (Criminal West) }

Rule {
  x, y, z,
  { (American x) (Weapon y) (Sells x y z) (Hostile z) }
  { (Criminal x) }
}

Rule {
  x
  { (Missile x) (Owns Nono x) }
  { (Sells West x Nono) }
}

Rule {
  x
  { (Missile x) }
  { (Weapon x) }
}

Rule {
  x
  { (Enemy x America) }
  { (Hostile x) }
}

```

By default all tuples in the graph are boolean-valued with default value true. In the above example all literals are actually true-valued. A rule $\{X, Y, \{ (p \ X \ Y) \ (q \ Y) \} \{ (q \ X) ! \} \}$ means $\forall XY \ p(X, Y)q(Y) \Rightarrow \neg q(X)$. If in the KB we only store true facts, this would 'delete' the fact $(q \ X) !$ from the KB (for some X).

As nodes of our graph can be of any type, we can represent predicates of any type, for instance $\{X, Y, \{ (p \ X \ Y) \ (q \ Y)=3 \} \{ (q \ X)=4 \} \}$ would let $p(X)$ be double-typed.

2.1 Methods

The most important methods are the following:

- Checking whether **two facts are equal**. Facts are grounded literals. Equality is simply checked by checking if all symbols (predicate or constant) in the tuples are equal. Optionally (by default), it is also checked if the fact values are equal.
- Checking whether **a fact equals a literal+substitution**. The literal is a tuple of symbols, some of which may be first order variables. All variables must be of the same scope (=declared in the same subgraph, in the same rule). The substitution is a mapping of these variables to root-level symbols (predicate or constant symbols). The method loops through the literal's symbols; whenever a symbol is in the substitution scope it replaces it by the substitution; then compares to the fact symbol. Optionally (by default) also the value is checked for equality.
- Check whether **a fact is directly true in a KB (or scope)** (without inference). This uses the graph connectivity to quickly find any KB-fact that shares the same symbols; then checks these for exact equality.
- Check whether **a literal+substitution is directly true in a KB** (without inference).
- Given a single literal with only ONE logic variable, and a KB of facts, **compute the domain** (=possible values of the variable) for the literal to be true. If the literal is negated the $D \leftarrow D \setminus d$, otherwise $D \leftarrow D \cup d$ if the d is the domain for true facts in the KB. [TODO: do this also for multi-variable literals]
- **Compute the set of possible substitutions for a conjunction of literals** (typically precondition of a rule) to be true in a KB.
- **Apply a set of 'effect literals'** (RHS of a rule): generate facts that are substituted literals

Given these methods, forward chaining, or forward simulation (for MCTS) is straightforward.

3 Optim

The optimization code contributes a nice way to represent (structured) optimization problems, and a few basic solvers for constrained and unconstrained, black-box, gradient-, and hessian-available problems.

3.1 Representing Optimization Problems

The data structures to represent optimization problems are

A standard unconstrained problem

$$\min_x f(x) \quad (1)$$

```
typedef std::function<double(arr& df, arr& Hf, const arr& x)> ScalarFunction;
```

The double return value is $f(x)$. The gradient df is returned only on request (for $\text{df} \neq \text{NoArr}$). The hessian Hf is returned only on request. If the caller requests df or Hf but the implementation cannot compute gradient or hessian, the implementation should `HALT`.

These can be implemented using a lambda expression or setting it equal to a C-function. See the examples.

A sum-of-squares problem (Gauss-Newton type)

$$\min_x y(x)^\top y(x) \quad (2)$$

```
typedef std::function<void(arr& y, arr& Jy, const arr& x)> VectorFunction;
```

where the vector y must always be returned. The Jacobian J_y is returned on request ($J_y \neq \text{NoArr}$).

A constrained problem (for vector valued functions f, y, g, h)

$$\min_x \sum_i f_i(x) + y(x)^\top y(x) \quad \text{s.t.} \quad g(x) \leq 0, \quad h(x) = 0 \quad (3)$$

Note that we can rewrite this as

$$\min_x \sum_{t \in F} \phi_t(x) + \sum_{t \in Y} \phi_t(x)^\top \phi_t(x) \quad \text{s.t.} \quad \forall_{t \in G} : \phi_t(x) \leq 0, \quad \forall_{t \in H} : \phi_t(x) = 0, \quad (4)$$

where the vector-valued *feature* function ϕ contains all f_i, y_i, g_i, h_i , and the disjoint partition $F \cup Y \cup G \cup H = \{1, \dots, T\}$ indicates whether the t -th feature contributes a scalar objective, sum-of-square objective, inequality constraint or equality constraint. We represent this as

```
enum TermType { noTT=0, fTT, sumOfSqrTT, ineqTT, eqTT };
typedef mlr::Array<TermType> TermTypeA;
struct ConstrainedProblem{
    virtual ~ConstrainedProblem() = default;
    virtual void phi(arr& phi, arr& J, arr& H, TermTypeA& tt, const arr& x) = 0;
};
```

Here, the returned `phi` is the feature vector and the returned `tt` indicates for every `phi`-entry its type (`fTT`, `sumOfSqrTT`, `ineqTT`, `eqTT`). The (on request) returned `J` is the Jacobian of `phi`. The (on request) returned `H` is the Hessian of the scalar features only, that is, the Hessian of $\sum_i f_i(x)$.

A **structured constrained problem**: Assume we have N decision variables $x_i \in \mathbb{R}^{d_i}$, each with its own dimensionality d_i . Assume we have J features $\phi_{j=1,\dots,J}$, but each feature ϕ_j depends on only a tuples $X_j \subseteq \{x_1, \dots, x_N\}$ of variables. We minimize

$$\min_{x_{1:N}} \sum_{j \in F} \phi_j(X_j) + \sum_{j \in Y} \phi_j(X_j)^\top \phi_j(X_j) \quad \text{s.t.} \quad \forall_{j \in G} : \phi_j(X_j) \leq 0, \quad \forall_{j \in H} : \phi_j(X_j) = 0. \quad (5)$$

```
struct GraphProblem {
    virtual void getStructure(uintA& variableDimensions, uintAA& featureVariables, TermTypeA& featureTypes) = 0;
    virtual void phi(arr& phi, arrA& J, arrA& H, const arr& x) = 0;
};
```

Here we decided to provide a method that first allows the optimizer to query the structure of the problem: return $N, d_{i=1,\dots,N}, J, X_{t=1,\dots,J}$, and $\text{tt}_{i=1,\dots,J}$. This allows the optimizer to setup its own data structures or so. Then, in each iteration the optimizer only queries $\text{phi}(\dots)$. This always returns the J -dimensional feature vector phi , which contains an f_i, y_i, g_i or h_i -value, depending on $\text{tt}(j)$. This $\text{phi}(j)$ may only depend on the decision variables X_j . On request it returns the gradient $J(j)$ of $\text{phi}(j)$ w.r.t. X_j . Note that the dimensionality of X_j may vary—therefore we return an array of gradients instead of a Jacobian. On request also a hessian $H(j)$ is returned for the scalar objectives (when $\text{tt}(j) == \text{fTT}$).

A **k-order Markov Optimization problem**. We have T decision variables $x_{1,\dots,T}$, each with potentially different dimensionality $d_{1,\dots,T}$. We have J features $\phi_{1,\dots,J}$, each of which may only depend on $k + 1$ consecutive variables $X_j = (x_{t_j-k}, \dots, x_{t_j})$, where t_j tells us which $k + 1$ -tuple feature ϕ_j depends on. We minimize again (5). For easier readability, this is equivalent to a problem of the form:

$$\min_{x_{1:T}} \sum_{t=1}^T y_t(x_{t-k:t})^\top y_t(x_{t-k:t}) \quad \text{s.t.} \quad \forall_t : g_t(x_{t-k:t}) \leq 0, \quad h_t(x_{t-k:t}) = 0, \quad (6)$$

where the feature with same $t_k = t$ have been collected in different vector-value functions y_t, g_t, h_t .

```
struct KOMO_Problem {
    virtual uint get_k() = 0;
    virtual void getStructure(uintA& variableDimensions, uintA& featureTimes, TermTypeA& featureTypes) = 0;
    virtual void phi(arr& phi, arrA& J, arrA& H, TermTypeA& tt, const arr& x) = 0;
};
```

Here, the structure function returns $N, d_{1,\dots,N}, J, t_j, \text{tt}(j)$.

4 KOMO

I do not introduce the KOMO concepts here. Read this¹ <http://ipvs.informatik.uni-stuttgart.de/mlr/papers/16-toussaint-Newton.pdf>!

¹ M. Toussaint: A tutorial on Newton methods for constrained trajectory optimization and relations to SLAM, Gaussian Process smoothing, optimal control, and probabilistic inference. In Geometric and Numerical Foundations of Movements, Springer, 2016.

The goal of the implementation is the separation between the code of optimizers and code to specify motion problems. The problem form (6) provides the abstraction for that interface. The optimization methods all assume the general form

$$\min_x f(x) \quad \text{s.t.} \quad g(x) \leq 0, \quad h(x) = 0 \quad (7)$$

of a non-linear constrained optimization problem, with the additional assumption that the (approximate) Hessian $\nabla^2 f(x)$ can be provided and is semi-pos-def. Therefore, the KOMO code essentially does the following

- Provide interfaces to define sets of k -order task spaces and costs/constraints in these task spaces at various time slices; which constitutes a MotionProblem. Such a MotionProblem definition is very semantic, referring to the kinematics of the robot.
- Abstracts and converts a MotionProblem definition into the general form (6) using a kinematics engine. The resulting MotionProblemFunction is not semantic anymore and provides the interface to the generic optimization code.
- Converts the problem definition (6) into the general forms (3) and (7) using appropriate matrix packings to exploit the chain structure of the problem. This code does not refer to any robotics or kinematics anymore.
- Applies various optimizers. This is generic code.

The code introduces specialized matrix packings to exploit the structure of J and to efficiently compute the banded matrix $J^\top J$. Note that the rows of J have at most $(k + 1)n$ non-zero elements since a row refers to exactly one task and depends only on one specific tuple (x_{t-k}, \dots, x_t) . Therefore, although J is generally a $D \times (T + 1)n$ matrix (with $D = \sum_t \dim(f_t)$), each row can be packed to store only $(k + 1)n$ non-zero elements. We introduced a *row-shifted* matrix packing representation for this. Using specialized methods to compute $J^\top J$ and $J^\top x$ for any vector x for the row-shifted packing, we can efficiently compute the banded Hessian and any other terms we need in Gauss-Newton methods.

4.1 Formal problem representation

The following definitions also document the API of the code.

KinematicEngine is a mapping $\Gamma : x \mapsto \Gamma(x)$ that maps a joint configuration to a data structure $\Gamma(x)$ which allows to efficiently evaluate task maps. Typically $\Gamma(x)$ stores the frames of all bodies/shapes/objects and collision information. More abstractly, $\Gamma(x)$ is any data structure that is sufficient to define the task maps below.

Note: In the code there is yet no abstraction KinematicEngine. Only one specific engine (KinematicWorld) is used. It would be straight-forward to introduce an abstraction for kinematic engines pin-pointing exactly their role for defining task maps.

TaskMap is a mapping $\phi : (\Gamma_{-k}, \dots, \Gamma_0) \mapsto (y, J)$ which gets $k + 1$ kinematic data structures as input and returns some vector $y \in \mathbb{R}^d$ and (on request) its Jacobian $J \in \mathbb{R}(d \times n)$.

Task is a tuple $c = (\phi, \varrho_{1:T}, y_{1:T}^*, \text{tt})$ where ϕ is a TaskMap and the parameters $\varrho_{1:T}, y_{1:T}^* \in \mathbb{R}^{T \times d}$ allow for an additional linear transformation in each time slice. Here, $d = \dim(\phi)$ is the dimensionality of the task map. This defines the transformed task map

$$\hat{\phi}_t(x_{t-k}, \dots, x_t) = \text{diag}(\varrho_t)(\phi(\Gamma(x_{t-k}), \dots, \Gamma(x_t)) - y_t^*), \quad (8)$$

which depending on $tt \in \{\text{fTT}, \text{sumOfSqrTT}, \text{ineqT}, \text{eqT}\}$ is interpreted as cost or constraint feature. Note that, in the cost case, $y_{1:T}^*$ has the semantics of a reference target for the task variable, and $\varrho_{1:T}^*$ of a precision. In the code, $\varrho_{1:T}, y_{1:T}^*$ may optionally be given as 1×1 , $1 \times T+1$, $d \times 1$, or $d \times T+1$ matrices—and are interpreted constant along the missing dimensions.

MotionProblem is a tuple $(T, \mathcal{C}, x_{-k+1:0})$ which gives the number of time steps, a list $\mathcal{C} = \{c_i\}$ of Tasks, and a *prefix* $x_{-k:-1} \in \mathbb{R}^{k \times n}$. The prefix allows to evaluate tasks also for time $t \leq k$, where the prefix defines the kinematic configurations $\Gamma(x_{-k+1}), \dots, \Gamma(x_0)$ at negative times. This defines the KOMO problem.

4.2 User Interfaces

4.2.1 Easy

For convenience there is a single high-level method to call the optimization, defined in

```
<Motion/komo.h>
    /// Return a trajectory that moves the endeffector to a desired target position
    arr moveTo(ors::KinematicWorld& world, //in initial state
               ors::Shape& endeff,         //endeffector to be moved
               ors::Shape& target,         //target shape
               byte whichAxesToAlign=0,    //bit coded options to align axes
               uint iterate=1);            //usually the optimization methods may be called just
                                           //once; multiple calls -> safety
```

The method returns an optimized joint space trajectory so that the endeff reaches the target. Optionally the optimizer additionally aligns some axes between the coordinate frames. This is just one typical use case; others would include constraining vector-alignments to zero (orthogonal) instead of +1 (parallel), or directly specifying quaternions, or using many other existing task maps. See expert interface.

This interface specifies the relevant coordinate frames by referring to Shapes. Shapes (`ors::Shape`) are rigidly attached to bodies (“links”) and usually represent a (convex) collision mesh/primitive. However, a Shape can also just be a marker frame (`ShapeType markerST=5`), in which case it is just a convenience to define reference frames attached to bodies. So, the best way to determine the geometric parameters of the endeffector and target (offsets, relative orientations etc) is by transforming the respective shape frames (`Shape::rel`).

The method uses implicit parameters (grabbed from `cfg` file or command line or default):

```
double posPrec = MT::getParameter<double>("KOMO/moveTo/precision", 1e3);
double colPrec = MT::getParameter<double>("KOMO/moveTo/collisionPrecision", -1e0);
double margin = MT::getParameter<double>("KOMO/moveTo/collisionMargin", .1);
double zeroVelPrec = MT::getParameter<double>("KOMO/moveTo/finalVelocityZeroPrecision", 1e1);
double alignPrec = MT::getParameter<double>("KOMO/moveTo/alignPrecision", 1e3);
```


4.2.2 Using a specs file

```
KOMO{
  T = 100
  duration = 5
}

Task sqrAccelerations{
  map={ type=qItself }
  order=2 # accelerations (default is 0)
  time=[0 1] # from start to end (default is [0 1])
  type=cost # squared costs (default is 'cost')
  scale=1 # factor of the map (default is [1])
  target=[0] # offset of the map (default is [0])
}

Example: Task finalHandPosition{
  map={ type=pos ref1=hand ref2=obj vec1=[0 0 .1] }
  time=[1 1] # only final
  type=equal # hard equality constraint
}

Task finalAlignmentPosition{
  map={ type=vecAlign ref1=hand vec1=[1 0 0] vec2=[0 1 0] }
  time=[1 1] # only final
  type=equal # hard equality constraint
  target=[1] # scalar product between vec1@hand and vec2@world shall be 1
}

Task collisions{
  map={ type=collisionIneq margin=0.05 }
  type=inEq # hard inequality constraint
}
```

4.2.3 Expert using the included kinematics engine

See the implementation of `moveTo`! This really is the core guide to build your own cost functions.

More generically, if the user would like to implement new TaskMaps or use some of the existing ones:

- The user can define new k -order task maps by instantiating the abstraction. There exist a number of predefined task maps. The specification of a task map usually has only a few parameters like “which endeffector shape(s) are you referring to”. Typically, a good convention is to define task maps in a way such that *zero* is a desired state or the constraint boundary, such as relative coordinates, alignments or orientation. (But that is not necessary, see the linear transformation below.)
- To define an optimization problem, the user creates a list of tasks, where each task is defined by a task map and parameters that define how the map is interpreted as a) a cost term or b) an inequality constraint. This interpretation allows: a linear transformation separately for each t (=setting a reference/target and precision); how maps imply a constraint. This interpretation has a significant number of parameters: for each time slice different targets/precisions could be defined.

4.2.4 Expert with own kinematics engine

The code needs a data structure $\Gamma(q_t)$ to represent the (kinematic) state q_t , where coordinate frames of all bodies/shapes/objects have been precomputed so that evaluation of task maps is fast. Currently this is `KinematicWorld`.

Users that prefer using the own kinematics engine can instantiate the abstraction. Note that the engine needs to fulfill two roles: it must have a `setJointState` method that also precomputes all frames of all bodies/shapes/objects. And it must be sufficient as argument of your task map instantiations.

4.2.5 Optimizers

The user can also only use the optimizers, directly instantiating the k -order Markov problem abstraction; or, yet a level below, directly instantiating the `ConstrainedProblem` abstraction. Examples are given in `examples/Optim/kOrderMarkov` and `examples/Optim/constrained`. Have a look at the specific implementations of the benchmark problems, esp. the `ParticleAroundWalls` problem.

4.2.6 Parameters & Reporting

Every run of the code generates a `MT.log` file, which tells about every parameter that was internally used. You can overwrite any of these parameters on command line or in an `MT.cfg` file.

Inspecting the cost report after an optimization is important. Currently, the code goes through the task list `C` and reports for each the costs associated to it. There are also methods to display the cost arising in the different tasks over time.

4.3 Potential Improvements

There is many places the code code be improved (beyond documenting it better):

- The `KinematicEngine` should be abstracted to allow for easier plugin of alternative engines.
- Our kinematics engine uses `SWIFT++` for proximity and penetration computation. The methods would profit enormously from better (faster, more accurate) proximity engines (signed distance functions, sphere-swept primitives).

4.4 Disclaimer

This document by no means aims to document all aspects of the code, esp. those relating to the used kinematics engine etc. It only tries to introduce to the concepts and design decisions behind the KOMO code.

More documentation of optimization and kinematics concepts used in the code can be drawn from my teaching lectures on Optimization and Robotics.

5 Control

5.1 Slow and fast control loop

There are two nested control loops:

In the slow loop ($\sim 50\text{Hz}$, non-strict, non-real-time) the controller has full access to the results of pre-computed optimizations, full models of the robots kinematics (dynamics?) and potentially delayed information on the robot (current pose, forces, contacts, etc). The slow loop may realize computationally complex things, e.g., operational space control, re-adaptation of a plan (phase adaptation, recalibration of task maps), model predictive control, online planning, etc.

The fast loop is 1kHz , strictly and real-time. It has direct access to the current robot state q (needs to compute \dot{q} from filtered differentiation of q) as well as the current readouts of the F/T sensors u_{ft} . **We constrain the fast controller to be a linear regulator in these**

observables and their integral:

$$e \leftarrow \gamma e + (f^* - J_{\text{ft}}^\dagger u_{\text{ft}}) \quad \text{or} \quad \dot{e} = (f^* - J_{\text{ft}}^\dagger u_{\text{ft}}) + (1 - \gamma)e \quad (9)$$

$$u = u_0 + k_p^{\text{base}} \cdot K_p(q^* - q) + k_d^{\text{base}} \cdot K_d(\dot{q}^* - \dot{q}) + K_I e. \quad (10)$$

This a (redundant) parameterization of a regulator linear in (q, \dot{q}, e) . We choose this parameterization because q^*, \dot{q}^*, f^* can be interpreted as “references”. But actually, we could just drop them (absorb them in u_0) without losing generality. In addition to this, the fast loop respects control limits by clipping $u \leftarrow \text{clip}(u, -u_{\text{max}}, u_{\text{max}}$ element-wise. u_{max} is a constant set in configuration files, not a fluent. [TODO: Additional mechanisms should also in the fast loop guarantee velocity and joint limits.]

The parameter vectors k_p^{base} and k_d^{base} are constants set in the PR2 configuration files. They are hand-tuned so that setting $K_p = K_d = I$ leads to acceptable (rather low gain) behavior. The \cdot denotes an element-wise product.

About the integral term: J_{ft}^\dagger allows us to linearly project the sensor signals to any other space in which we have a target f^* and integrate the error.

$K_p, K_d, K_I, J_{\text{ft}}^\dagger$ are arbitrary matrices; u_0 an arbitrary control bias. Therefore, the **control mode** of the fast loop is determined by the tuple

$$M = (q^*, \dot{q}^*, f^*, u_0, K_p, K_d, K_I, J_{\text{ft}}^\dagger, \gamma). \quad (11)$$

This is the message that the slow loop needs to send to the fast loop – the slow loop can change the control mode at any time.

Inversely, the fast loop passes the message

$$(q, \dot{q}, f, u) \quad (12)$$

to the slow loop, giving it information on the true current state (q, \dot{q}) , sensor readings f , and computed controls u .

The core question therefore is how the slow loop computes the message M to realize the desired control behaviors. The list of basic desired control behaviors is:

1. Follow a pre-computed trajectory $(q_{0:T}, \tau)$, where τ is the time resolution
2. Follow the position-reference that is online computed by a operational space (or inverse kinematics) controller; the K_p should such that P-gains can be set/added/removed along *endeffector* spaces rather than only uniformly configuration space
3. Establish a contact
4. Stabilize a contact force
5. Limit F/T (to avoid breaking a handle)
6. Sliding (moving tangentially) on a surface (or along a DOF like an opening door) which is perceived via the F/T signal

5.2 Operational Space Control: Computing gains by projecting operational space gains

The appendix B derived the necessary equations in all generality. In practise, it is sufficient to modify the K_p only, using the Jacobian of the desired task space. In equation (59) we have

$$\bar{K}_p = A^{-1} J^\top C K_p J, \quad A = H + J^\top C J, \quad (13)$$

where we assumed $M = \mathbf{I}$ (quasi-dynamic model) and no other tasks. Further, assuming $C = c$ and $K_p = k$ are scalars we have

$$\bar{K}_p = k(H/c + J^\top J)^{-1} J^\top J. \quad (14)$$

I actually tested just $kJ^\top J$.

TODO: Let `FeedbackController` really compute these projected PD behaviors, instead of only q^*, \dot{q}^* ! Then all of this is automatic!

5.3 Controlling the F/T signal—the *sensed* force

5.3.1 Preliminaries: Understanding force transmission

The following law of force propagation is well known,

$$u = J^\top f \quad (15)$$

where f is a force in the endeffector (e.g., the negative of its gravity load), J the position Jacobian of the endeffector, and u are the torques “perceived” in each of the joints due to the force f . This law is correct only under the assumption that nothing moves. Inversely, this law is typically used to compensate forces: Assume you have a load on an endeffector, gravity pulls it down. The gravity force pulling the load down propagates to torques u in each joint – if you want to compensate this torque the motors need to create the reactive torque.

The same also holds for force-torque $f \in \mathbb{R}^6$, where the Jacobian is the stacking of the position and the axial Jacobian.

Typically, f is lower-dimensional than u . So, actually, there should be many u that generate a desired f^* ? What is the optimal one? Well, assume $f^* = 0$ for a moment. Then, any choice of u will accelerate the robot (assuming gravity compensation). The only choice to generate $f^* = 0$ and not to accelerate the robot is $u = 0$. Equally, the only choice to generate any f^* without accelerating the robot is $u = J^\top f^*$.

When we include system dynamics in the equation, we have the general

$$u = M\ddot{q} + h + J^\top f. \quad (16)$$

where M (the inertia matrix) and h (the coreolis and gravity forces) depend on (q, \dot{q}) . One way to read this equation is: the torques you “feel” in the joints are the reactive torques of the robot’s inertia (that derive the acceleration) plus the torque you feel from the endeff force f .

5.3.2 Controlling the direct F/T signal with a fixated endeff

Consider the following exercise: Fix the endeffector rigidly, e.g. to a table with a clamp (Schraubzwinge). Write a controller that generates any desired f^* in the F/T sensor with the least effort, and stably, and staying close to a desired homing posture.

If we unrealistically assume that our model is correct then the solution simply is (16); for $\ddot{q} = 0$ and a gravity-compensated robot just (15); where

$$J = J_{\text{ft}}, \quad (17)$$

which is the position and axial Jacobian of the F/T sensor w.r.t. q .

However, this equation **does not use any F/T sensor feedback** to generate the desired F/T signal. This cannot work well in practise. We can resolve this with an I-controller on the F/T signal error.

$$e = \int_t dt [f^* - f] \quad (18)$$

$$u = J^\top \alpha e. \quad (19)$$

The α here has the meaning of an exponential decay of the signal error—which we can show assuming the perfect model. Under perfect model assumption, the F/T sensor measures

$$f = J^\dagger u, \quad J^\dagger J^\top \equiv \mathbf{I}, \quad J^\dagger = (J J^\top)^{-1} J \quad (20)$$

$$\text{Note: } J^\dagger u = J^\dagger J^\top f = (J J^\top)^{-1} J J^\top f = f \quad (21)$$

$$(22)$$

Note that $J J^\top$ is a $d \times d$ -matrix and invertible and J^\dagger the appropriate left-pseudo-inverse of J^\top . Inserting this perfect-model measurement in the control law (18) we get

$$\dot{e} = f^* - J^\dagger u \quad (23)$$

$$\dot{u} = J^\top \alpha (f^* - J^\dagger u) = \alpha J^\top f^* - \alpha \underbrace{J^\top J^\dagger}_{= \mathbf{I}} u = \alpha (J^\top f^* - u). \quad (24)$$

Here, $J^\top J^\dagger = J^\top (J J^\top)^{-1} J^\top$ is actually the projection that projects any joint torques u into the space that directly relates to endeffector forces and not to accelerations. However, if the u was chosen by some law $J^\top f$, then u will always lie within this projection (will never lead to accelerations of the robot), and therefore it actually is the identity matrix.

Now, the above states that $\dot{u} = \alpha (J^\top f^* - u)$, which says that u exponentially approaches the perfect-model correct torque $J^\top f^*$, which a decay rate α . Therefore, α can be considered a decay rate.

Open: What if we have a \ddot{q} as well? Two possibilities: 1) Reiterate the above reasoning with \ddot{q} . 2) Just add the signals.

5.3.3 Control the indirectly sensed contact force of endeff

Exercise: We have the F/T sensor, but attached to it a hand and a contact point with some relative transformation to the F/T sensor. This point is in contact with a table. What we want to control is the force between point and table, which is just a 1D thing.

This is best addressed by thinking of the F/T sensor as if it was a 6D joint (like a ball joint). If we have a force f at some endeffector then we “feel” this force in all joints of the robot as $u = J^\top f$. This includes the F/T sensor joints! So the Jacobian of the endeffector variable (be it 1D or 3D) w.r.t. the sensor pseudo-ball-joint exactly gives the measurement equation. Let’s denote this Jacobian as

$$J_{ft} \in \mathbb{R}^{d \times 6}, \quad (25)$$

where $d = 1$ if it is only the distance to the table, or $d = 3$ if it is all forces. Further, let’s denote by

$$J \in \mathbb{R}^{d \times n} \quad (26)$$

the Jacobian w.r.t. all the real robot joints.

As above, the *perceived* endeffector force (this time perceived by the F/T sensor) is

$$u_{ft} = J_{ft}^\top f \Rightarrow f = J_{ft}^\dagger u_{ft} , \quad (27)$$

where $u_{ft} \in \mathbb{R}^6$ is the F/T signal. Again we may use an I-controller to correct for the error between desired endeffector force f^* and perceived one:

$$\dot{e} = f^* - f = f^* - J_{ft}^\dagger u_{ft} \quad (28)$$

$$u = J^\top \alpha e . \quad (29)$$

Note that the last equation generates joint torques proportional to the normal endeffector Jacobian J because e is an error in endeffector force space (not F/T signal space).

This fits to our controller setup by

$$J_{ft}^\dagger \leftarrow J_{ft}^\dagger , \quad f^* \leftarrow f^* , \quad \gamma \leftarrow 1 , \quad K_I \leftarrow \alpha J^\top . \quad (30)$$

When force control is turned off, we need to remember to set $\gamma = 0, e = 0$ to ensure that next time it is turned on again it doesn't blow.

Open: What happens for $\gamma < 1$? Is this equivalent to $\alpha < 1$? Perhaps not. ($\sum_{t=0}^{\infty} \gamma^t = \frac{1}{1-\gamma}$)

5.3.4 q -control under force constraints

Assume we have a P(I)D controller on q —typically a PID in some task space that has been projected to joint space. We would like to execute that desired reference behavior but subject to constraints on the sensed endeffector force

$$f_{lo} \leq J_{ft}^\dagger f \leq f_{hi} . \quad (31)$$

These are $2d$ constraints.

As with lagrange parameters, we can simply activate the constraints when violated: When one of the components violates the constraint, control the force to be exactly $f_{lo|hi}$. For the latter, use the f -error-integral method as above. This should eventually have higher priority to any other gains (keep other I-gains limited!).

5.3.5 I-gains on position?

5.4 Technical Details and Issues

5.4.1 Ctrl-Message documentation

One message type for setting the control mode AND feedback from the controller.

Setting the control model: $(q^*, \dot{q}^*, f^*, u_0, K_p, K_d, K_f)$. Can be set any time.

Feedback from the controller: (q, \dot{q}, f, u) . Published with 1kHz.

5.4.2 Filtering of the differentiation of q

[Peter: please fill in]

5.5 Enforcing control, velocity, joint and force limits

Enforcing control limits is really simple: Just clip the computed u .

Enforcing velocity limits turned out difficult: The velocity signal is so noisy, a direct feedback coupling was bad. Also, the IF-case of velocity-limit-violation turned off and on quickly and introduced even more noise (rattling motors...)

I have ignored limits totally so far – should be handled (as collisions) in the slow loop.

FORCE LIMITS! Not idea how to handle this.

Maybe a route to a more principled approach to all of these: Take the Augmented Lagrangian way to handle constraints as a template: First associate only a soft squared penalty with margin penetration. Then compute/update the respective dual parameters that push you out of the margin.

.1 Reference: General OPEN-LOOP ideal contact force controller

.1.1 General case

This controller is sort-of open loop! It does not take into account any F/T feedback. What it receives as a specification is \ddot{y}^* (desired endeff accel) and λ^* (desired contact force); as well as models of the maps ϕ, J_ϕ, g, J_g . We discuss below how this can properly be made a feedback regulator. We write the problem as a general constraint problem

$$\min_{u, \ddot{q}, \lambda} \|u - a\|_H^2 \quad (32)$$

$$\text{s.t. } u = M\ddot{q} + h + J_g^\top \lambda \quad (33)$$

$$J_\phi \ddot{q} = c \quad (34)$$

$$\lambda = \lambda^* \quad (35)$$

$$J_g \ddot{q} = b \quad (36)$$

Notes:

- The role of a becomes clearer when we treat the blue constraint below
- The 2nd constraint relates to an arbitrary task map $\phi : q \mapsto y$ with Jacobian J_ϕ , $c = \ddot{y}^* - \dot{J}_\phi \dot{q}$ and some desired task space acceleration \ddot{y}^* .
- We have a set of functions $g : q \rightarrow \mathbb{R}^m$ with Jacobian J_g which play the role of inequality constraints
- The 3rd constraint captures desired contact forces with the constraints
- The blue constraint expresses that a) assuming the contact is active there must not be acceleration w.r.t. g or b) if the contact is not active we might want to control the acceleration towards it (to make it active). (This constrains the dynamics. Without this constraint the dynamics could assume that the contact forces λ are generated while the endeff is moving (e.g., by strange external forces). This constraint makes it consistent.)

To derive closed form solutions, each of these equality constraints can be handled in two ways: relax it to become a squared penalty (and then potentially taking the infinite precision limit); or resolve it.

We resolve the 1st and 3rd constraint, and relax the 2nd and 4th to later take the limit. The solution is

$$\mathbf{h} := h + J_g^\top \lambda^* \quad (37)$$

$$f(\ddot{q}) = \|M\ddot{q} - (a - \mathbf{h})\|_H^2 + \|J_g \ddot{q} - b\|_B^2 + \|J_\phi \ddot{q} - c\|_C^2 \quad (38)$$

$$\ddot{q}^* = (M^\top H M + J_g^\top B J_g + J_\phi^\top C J_\phi)^{-1} [M^\top H (a - \mathbf{h}) + J_g^\top B b + J_\phi^\top C c] \quad (39)$$

The limit $B \rightarrow \infty$:

$$\ddot{q}^* = (X + J_g^\top B J_g)^{-1} [J_g^\top B b + x] \quad (40)$$

$$= J_{gXB}^\# b + (\mathbf{I} - J_{gXB}^\# J_g) (X^{-1} x) \quad (41)$$

And note that $(X^{-1}x)$ is the solution to only having the other terms. Given \ddot{q}^* , the optimal control is computed as $u = M\ddot{q} + h + J_g^\top \lambda^*$. We still did not take the limit of the C -term (endeffector position control). We could use the hierarchical limit case.

.2 Reference: Pullback of operational space linear controllers

The above assumes that at any instance in time we want a certain task-space acceleration \ddot{y}^* and translates this to an optimal joint control in that instant in time. If we want to implement a certain feedback behavior in the task space, that is, we have a desired feedback control law $\pi : y, \dot{y} \mapsto \ddot{y}$, we can evaluate π at every point in time and project to operational space control.

$$\ddot{y} = \ddot{\phi}(q) = \ddot{(Jq)} = (\dot{J}q + J\dot{q}) = 2\dot{J}\dot{q} + J\ddot{q} \quad (42)$$

$$\ddot{y}^* = K_p y + K_d \dot{y} + k \quad (43)$$

$$J\ddot{q} \stackrel{!}{=} c = \ddot{y}^* - 2\dot{J}\dot{q} = K_p y + K_d \dot{y} + k - 2\dot{J}\dot{q} \quad (44)$$

$$\approx K_p (J(q - q_0) + \phi(q_0)) + K_d J\dot{q} + k \quad (45)$$

$$= K_p Jq + K_d J\dot{q} + k', \quad k' = k + K_p(\phi(q_0) - Jq_0) \quad (46)$$

$$\ddot{q}^* = A^{-1}[\dots + J^\top C c] = A^{-1}[\dots + J^\top C (K_p Jq + K_d J\dot{q} + k')] \quad (47)$$

$$= A^{-1}[\dots] + A^{-1} J^\top C K_p Jq + A^{-1} J^\top C K_d J\dot{q} + A^{-1} J^\top C k' \quad (48)$$

$$= \bar{K}_p q + \bar{K}_d \dot{q} + \bar{k}, \quad \bar{k} = A^{-1}[\dots] + A^{-1} J^\top C k', \quad \bar{K}_p = A^{-1} J^\top C K_p J, \quad \bar{K}_d = A^{-1} J^\top C K_d J \quad (49)$$

.3 How to make this FEEDBACK?

W.r.t. y (endeff pos) it is clear how to make this feedback: We can impose a PD behavior on the endeffector

$$\ddot{y}^* = k_p(y^* - y) + k_d(\dot{y}^* - \dot{y})$$

and send this desired endeff accel to the general controller.

What about λ^* ??

.4 What do we want?

desired task space acceleration law

$$\ddot{y}^* = \ddot{y}^{\text{ref}} + K_p(y^{\text{ref}} - y) + K_d(\dot{y}^{\text{ref}} - \dot{y}) + K_{Ip} \int (y^{\text{ref}} - y) + K_{Id} \int (\dot{y}^{\text{ref}} - \dot{y}) \quad (50)$$

That defines a desired *acceleration*. But if the system was precise in enforcing this acceleration it would be non-compliant. Note: strictly speaking, if this law says $\ddot{y}^* = 0$ (e.g., because K_p and K_d are zero), then this is a strict (non-compliant) statement.

precision/compliance Given a desired \ddot{y}^* , the precision along some dimensions may not be that important. We may capture this with the precision (or compliance) matrix C . As a convention, let the $\text{eig}(C) \in [0, 1]$, and an eigenvalue of 1 states full precision, while an eigenvalue of 0 states full compliance.

This implies an objective term

$$\|J_\phi \ddot{q} - \ddot{y}^*\|_C^2$$

Null cost reference Typically one defines control costs $\|u\|_H^2 = \|M\ddot{q} + F\|_H^2$. However, this becomes semantically tricky. When defining what is ‘desired’ I propose to stay on the level of accelerations. So we have a desired (‘null’) acceleration law

$$\ddot{q}_0^* = \ddot{q}_0^{\text{ref}} + K_p^q(q_0^{\text{ref}} - q) + K_d^q(\dot{q}_0^{\text{ref}} - \dot{q}) \quad (51)$$

and consider costs

$$\|\ddot{q} - \ddot{q}_0^*\|_H^2.$$

Note that, by appropriate choices of parameters, the typical control cost can be mimicked. However, the semantics is somewhat different. For instance, setting K_d^q and \dot{q}_0^{ref} implies that we want to damp motion, and choosing controls u that implement this damping are at *null costs*. Equally, for non-zero F and $\ddot{q}_0^* = 0$, applying controls that ensure zero acceleration are at *null costs*. That’s what I call it *null cost reference*. This is rather different to generally penalize $\|u\|_H^2$, which would imply costs for any controls u , even if they just implement counteracting F of generating the null reference.

Optimal acceleration law We compute the optimal acceleration \ddot{q}^* in its 1st order Taylor approximation w.r.t. q and \dot{q} :

$$\ddot{y}^* = \ddot{y}^{\text{ref}} + K_p(y^{\text{ref}} - y) + K_d(\dot{y}^{\text{ref}} - \dot{y}) \quad (52)$$

$$\approx \ddot{y}^{\text{ref}} + K_p y^{\text{ref}} - K_p(J(q - q_0) + y_0) + K_d \dot{y}^{\text{ref}} - K_d J \dot{q} \quad (53)$$

$$= k - K_p J q - K_d J \dot{q}, \quad k = \ddot{y}^{\text{ref}} + K_p y^{\text{ref}} + K_d \dot{y}^{\text{ref}} + K_p(J q_0 - y_0) \quad (54)$$

$$\ddot{q}_0^* = k^q - K_p^q q - K_d^q \dot{q}, \quad k^q = \ddot{q}_0^{\text{ref}} + K_p^q q_0^{\text{ref}} + K_d^q \dot{q}_0^{\text{ref}} \quad (55)$$

$$\ddot{q}^* = \underset{\ddot{q}}{\text{argmin}} \|\ddot{q} - \ddot{q}_0^*\|_H^2 + \|J\ddot{q} - \ddot{y}^*\|_C^2 \quad (56)$$

$$= (H + J^\top C J)^{-1} [H \ddot{q}_0^* + J^\top C \ddot{y}^*] \quad (57)$$

$$\approx \bar{k} - \bar{K}_p q - \bar{K}_d \dot{q} \quad (58)$$

$$= \bar{K}_p(q_{\text{ref}} - q) - \bar{K}_d \dot{q}, \quad q_{\text{ref}} = \bar{K}_p^{-1} \bar{k} \quad (59)$$

$$\bar{K}_p = (H + J^\top C J)^{-1} [H K_p^q + J^\top C K_p J] \quad (60)$$

$$\bar{K}_d = (H + J^\top C J)^{-1} [H K_d^q + J^\top C K_d J] \quad (61)$$

$$\bar{k} = (H + J^\top C J)^{-1} [H k^q + J^\top C k] \quad (62)$$

Transfer to controls Given an optimal acceleration law and the system dynamics, we choose:

$$\ddot{q}^* = \bar{K}_p q + \bar{K}_d \dot{q} + \bar{k} \quad (63)$$

$$u^* = M \ddot{q}^* + F \quad (64)$$

$$= M \bar{K}_p q + M \bar{K}_d \dot{q} + M \bar{k} + F \quad (65)$$

Note again, only here, the system dynamics enter. The specification of the optimal acceleration law is independent of the dynamics. (Unless H and \ddot{q}^{ref} are chosen in relation to M and F to mimic typical control costs—but we explicitly avoid this to make desired system behavior somewhat less dependent on (possibly inaccurate) dynamics models).

Error correction of system dynamics The optimal acceleration law gives an explicit desired acceleration. We may estimate the control error by a low pass filter on the acceleration errors. There are two options:

Let $\langle \ddot{q}^* \rangle$ be a low pass filter of the desired accelerations $\ddot{q}^*(q, \dot{q})$; and $\langle \ddot{q} \rangle$ a low pass of the actual true accelerations. We may define $g = \langle \ddot{q} \rangle - \langle \ddot{q}^* \rangle$ and control

$$u = M(\ddot{q}^* + g) + F.$$

Or we may assume system dynamics

$$u = M\ddot{q} + F + g \tag{66}$$

for some unknown and variable (slowly changing) $g \in \mathbb{R}^n$ which reflects constant loads on the joints. We may estimate g as a low-pass filter,

$$g = \langle u - M\ddot{q} - F \rangle_{\text{low pass}} = \langle u \rangle - M \langle \ddot{q} \rangle - F \tag{67}$$

$$= \langle M\ddot{q}^* + F + g \rangle - M \langle \ddot{q} \rangle - F \tag{68}$$

$$= M [\langle \ddot{q}^* + g \rangle - \langle \ddot{q} \rangle] \tag{69}$$

which is puzzlingly different to the above. (M is obvious, but the rest?)

Compliant error correction of system dynamics The above describes a scheme that corrects any errors from the desired acceleration. However, in the case of contact, and desired compliance, we do not *want* to enforce exact reference following along certain dimensions. E.g., in the case of a contact we systematically do not accelerate towards, leading to a systematic error in the system equations, an adaptation of g , and perhaps divergence.

Error correction on task space level Let $\langle \ddot{y}^* \rangle$ be a low pass filter of the desired task space accelerations and $\langle \ddot{y} \rangle$ a low pass of the actual true accelerations. We may define $g = \langle \ddot{y} \rangle - \langle \ddot{y}^* \rangle$ and add g to the \ddot{y}^{ref} .

Alternatively, we may define g as integral error in the task space and add it to \ddot{y}^{ref} .

In both cases, g adds to k , showing that it adds to \ddot{q}^* as $(H + J^T C J)^{-1} J^T C g$.

The matrix C controls compliance.

Measured-force limits

Limit Energy