

# Learn to Throw a Ball

Sven, Luis, Heiko

July 13, 2016

Robot tries to hit a **target**, specified by an *Alvar-marker*.

- Controller:

$$\mu(W, t, s) = W\phi_t(s)$$

- Features:

$$\phi_t(s) = \left( t, \sin\left(\frac{t\pi}{T}\right), x_{eff}^{rot}, 1 \right)^T$$

- Robot keeps on sending *velocities* until **either**:
  - $T - 15$  commands have been send, **or**
  - $x$ -rot. of the gripper,  $x_{eff}^{rot}$ , is above a threshold (0.005)
- Distance between ball and target is measured placing a second Alvar-marker where the ball landed.

# Algorithm: Greedy Policy Search

```
Init  $W$  randomly with  $w_{ij} \sim \mathcal{U}(0, 1)$ ,  $r_{max} \leftarrow -10.000$   
while not converged do  
     $i \leftarrow rand(1, 2, 3)$  ▷ Use pr. sweeping  
    Add noise  $\varepsilon \sim \mathcal{N}(0, \sigma)$  to  $w_i$   
    for  $t \in T_i$  do  
        Get velocities for each joint:  $v_t \leftarrow \mu(W, \phi(s))$   
        Apply current velocity vector  $v_t$   
    end for  
    Get reward  $r = -sq\_dist(target, \hat{x})$  ▷  $\hat{x}$ : measured position  
    if  $r > r_{max}$  then  
         $W \leftarrow W + \varepsilon$   
        Reset  $\sigma$   
    else  
         $\sigma \leftarrow \sigma \cdot 1.1$   
    end if  
end while
```

During execution: collect *Data*  $\mathcal{D} = \{W_{11}^i, \dots, W_{34}^i, x_i, y_i, R_i\}_{i=1}^N$ .

Idea 1: Extract relevant data points around point of interest (*kNN*)

Use *Linear Regression* on the new data set to estimate *gradient*.

Idea 2: Use *supervised learning* to estimate *model* (on either  $W \mapsto R$  or  $W \mapsto (x, y)$ ).